

Vishal Vaka

+1-872-258-6648 | vishalvaka@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

MS CS graduate & end-to-end ML engineer who transforms research into production: curates data, trains Bayesian-tuned deep-learning / LLM pipelines in PyTorch + Hugging Face, and ships via cloud-native MLOps and CI/CD. Delivered real-time optimization engines, scalable recommender prototypes, and analytics dashboards that bring high-impact AI to non-experts.

EDUCATION






University of Illinois at Chicago | Master of Science in Computer Science | GPA - 3.78/4.0 Aug 2023 – May 2025
Coursework: Advanced Machine Learning, Artificial Intelligence, Parallel Processing, Natural Language Processing

Osmania University, Hyderabad | Bachelor of Technology in Computer Science and Engineering Jul 2017 – May 2021
Coursework: Machine Learning, DBMS, BigData, Python App Programming, Data Science, Data Analytics

EXPERIENCE

- **Machine Learning Researcher | PyTorch, BoTorch, GPyTorch, Lab Streaming Layer, GitHub Actions, Jupyter** Feb 2024 - Aug 2024
Chicago, USA
University Of Illinois at Chicago, Rehab Robotics Lab
 - Built real-time LSL to BoTorch pipeline streaming gait / metabolic data; GPU-optimised Bayesian loops ran 25 % faster.
 - Designed multi-objective Rank-Weighted GP ensemble with Chebyshev-EHVI, doubling Pareto coverage and halving convergence iterations.
 - Refactored RGPE into pip package; added 2-D benchmarks, Matplotlib diagnostics, Jupyter workflows, automated hyper-parameter sweeps.
 - Integrated PyTest and GitHub Actions CI/CD, lifting nightly build stability to 100 % across more than 70 commits.
 - Resolved clock-sync preprocessing bugs and wrote dev docs, enabling reproducible multi-GPU experiments for incoming researchers.
- **Business Technology Analyst | Salesforce, Apex, Jenkins, GitHub, Jira, Confluence, Apttus CPQ, CI/CD** Sep 2021– Aug 2023
Hyderabad, India
Deloitte USI
 - Automated CPQ deployments via Jenkins-GitHub pipeline; Apex unit tests and quality gates push code from QA to PROD.
 - Generated multilingual quotes by crafting Apttus X-Author templates and email microservice, slashing global-entity sales cycles 30%.
 - Implemented Apex price-list engine with country-specific rules; cut Belgium staging overrides 40% and improved pricing accuracy.
 - Achieved 95% release stability through exhaustive test classes, telemetry logging, and Jira-Confluence traceable CI documentation.

PROJECTS

- **SoloRAG – Self-Hosted RAG Stack | Python, FastAPI, Sentence-Transformers, FAISS, Ollama, Docker, Prometheus** 
 - Shipped one-command Docker RAG stack; offline fully deployment starts in 10 minutes using under 4 GB memory locally.
 - Built 15 MB FAISS index; retrieval boosts Exact-Match 27 points, achieving EM 48 and F1 64.7 on Stripe FAQs.
 - Added custom Prometheus middleware, 22 PyTests, 95 % coverage; GPU compose override slashes inference latency 4x without API changes.
 - CI/CD with GitHub Actions: tests, linting, Docker build ensure reliable, portable runs across Linux GPUs and CPUs.
- **Biomedical NER Comparison Suite | Python, PyTorch, HuggingFace Transformers, BioBERT, LoRA, CRF, Gradio, YAML** 
 - Benchmarked BioBERT fine-tuning vs LoRA, CRF, adapter fusion; reached 0.88 F1 using 98 % fewer parameters.
 - YAML runner executes triple-seed sweeps, logs latency + VRAM, finishes reproducible training in 15 min on a T4.
 - Published HuggingFace Spaces Gradio app—live token-level NER predictions with F1, runtime, memory dashboards.
 - Structured src package, PEP-8 lint, pinned requirements; MIT-licensed repo is CI/CD-ready for production ML workflows.
- **NLP Sentiment-&-Coherence Benchmark Toolkit | scikit-learn, Gensim Word2Vec, spaCy, Stanford CoreNLP, TF-IDF, SVM, Git LFS** 
 - Benchmarked TF-IDF, Word2Vec, handcrafted linguistic features across SVM, Logistic, MLP; macro-F1 rose from 0.57 to 0.71.
 - Added spaCy POS-diversity and conjunction-count extractors, boosting stylistic recall nine percent on social-sentiment data.
 - Integrated CoreNLP coherence metrics (TTR, LSA, content overlap) enabling real-time dialogue-quality dashboards for human-in-loop review.
 - Versioned models with Git LFS; CLI offers live classification, YAML sweeps, PEP-8 lint, CI-tested reproducible workflows.
- **Distributed Mandelbrot Fractal Renderer | C++, MPI, OpenMP, Complex Math, PBS, Linux** 
 - Built hybrid MPI + OpenMP Mandelbrot renderer; 188x faster than serial on 8-node, 64-core cluster.
 - Designed row-block decomposition and parallel MPI_File_write_at, removing gather bottleneck, sustaining 93 % strong-scaling efficiency.
 - Optimised complex-number loop with SIMD pragmas and cache-friendly tiling, boosting per-core FLOPS.
 - Automated reproducible HPC runs via Makefile, PBS scripts, and Python performance plots.
- **Distributed 2-D FFT Convolution Accelerator | C, MPI, OpenMP, FFT, PBS/SLURM, HPC** 
 - Parallelised 512 × 512 2-D FFT with MPI Scatter/Gather and OpenMP, achieving 7.8 × speedup over serial baseline.
 - Implemented communicator splitting and custom Scatterv/Gatherv, sustaining 92 % strong-scaling efficiency across 32 CPU cores.
 - Optimised cache-friendly transposes and SIMD complex-multiply kernels, maximising FLOPS for vision and transformer-attention convolutions.
 - Packaged Makefile and PBS scripts for reproducible SLURM runs, blueprinting distributed gradient-sync workflows for ML.

SKILLS

- **Programming Languages:** Python, C, C++, Java, Apex, JavaScript, HTML, CSS
- **ML/DL & Vector Retrieval:** PyTorch, TensorFlow, Keras, scikit-learn, HuggingFace Transformers, BoTorch, GPyTorch, FAISS, LoRA, CRF
- **Data Processing & Visualization:** NumPy, Pandas, Matplotlib, Seaborn, NLTK, spaCy, OpenCV, YAML
- **AI/ML:** LLMs, CNN, RNN, LSTM, Gaussian Processes, Bayesian Optimisation, Siamese Networks, Contrastive Learning, TF-IDF, Word2Vec, Sentiment Analysis, NER
- **Data Science:** Data Mining, Text Mining, Topic Modelling (LDA, LSA), Feature Engineering, Hyper-parameter Sweeps, Reproducible Pipelines
- **Natural Language Processing:** Sarcasm Detection, Named Entity Recognition, Semantic Analysis, Coherence Scoring
- **Tools & Platforms:** AWS SageMaker, Git, GitHub, Jenkins, GitHub Actions, Git LFS, Docker, Gradio, HuggingFace Spaces, Lab Streaming Layer, MPI, OpenMP, CUDA, SLURM/PBS, Salesforce (CPQ, Flows), Jira, Confluence

AFFILIATIONS

- **Vice Chair, Decent Work & Economic Growth | Hyderabad Youth Assembly (Season IX) – Street Cause - Hyderabad** Sep 2019 – Feb 2020
 - Promoted from Delegate to Vice Chair; mentored 20+ delegates executing SDG 8 projects.
 - Co-led orphanage and government-school outreach; installed equipment, taught environmental lessons, impacting 200+ students.
 - Designed arts-and-crafts micro-enterprise workshop; empowered 30 orphans to create and sell products for income.
 - Managed budget, vendors, and reports; ensured 100 % on-time delivery and transparent fund utilisation.